

```
In [25]: # importing useful libraries

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

Loading data separately for each file

```
In [26]: trans_df = pd.read_csv("Transactions.csv")
```

```
In [27]: product_df = pd.read_csv("Products.csv")
```

```
In [28]: custm_df = pd.read_csv("Customers.csv")
```

```
In [29]: trans_df.head()
```

```
Out[29]:
```

	TransactionID	CustomerID	ProductID	TransactionDate	Quantity	TotalValue	Price
0	T00001	C0199	P067	2024-08-25 12:38:23	1	300.68	300.68
1	T00112	C0146	P067	2024-05-27 22:23:54	1	300.68	300.68
2	T00166	C0127	P067	2024-04-25 07:38:55	1	300.68	300.68
3	T00272	C0087	P067	2024-03-26 22:55:37	2	601.36	300.68
4	T00363	C0070	P067	2024-03-21 15:10:10	3	902.04	300.68

checking for shape, null values, duplicates values for each file.

```
In [30]: trans_df.shape
```

```
Out[30]: (1000, 7)
```

```
In [31]: trans_df.isnull().sum()
```

```
Out[31]: TransactionID      0
         CustomerID        0
         ProductID         0
         TransactionDate    0
         Quantity           0
         TotalValue         0
         Price              0
         dtype: int64
```

```
In [32]: trans_df.duplicated().sum()
```

```
Out[32]: 0
```

```
In [33]: product_df.head()
```

```
Out[33]:
```

	ProductID	ProductName	Category	Price
0	P001	ActiveWear Biography	Books	169.30
1	P002	ActiveWear Smartwatch	Electronics	346.30
2	P003	ComfortLiving Biography	Books	44.12
3	P004	BookWorld Rug	Home Decor	95.69
4	P005	TechPro T-Shirt	Clothing	429.31

```
In [34]: product_df.isnull().sum()
```

```
Out[34]: ProductID      0
         ProductName    0
         Category       0
         Price          0
         dtype: int64
```

```
In [35]: product_df.duplicated().sum()
```

```
Out[35]: 0
```

```
In [36]: product_df.shape
```

```
Out[36]: (100, 4)
```

```
In [37]: custm_df.head()
```

```
Out[37]:
```

	CustomerID	CustomerName	Region	SignupDate
0	C0001	Lawrence Carroll	South America	2022-07-10
1	C0002	Elizabeth Lutz	Asia	2022-02-13
2	C0003	Michael Rivera	South America	2024-03-07
3	C0004	Kathleen Rodriguez	South America	2022-10-09
4	C0005	Laura Weber	Asia	2022-08-15

```
In [38]: custm_df.shape
```

```
Out[38]: (200, 4)
```

```
In [39]: custm_df.isnull().sum()
```

```
Out[39]: CustomerID      0
CustomerName    0
Region          0
SignupDate      0
dtype: int64
```

```
In [40]: custm_df.duplicated().sum()
```

```
Out[40]: 0
```

```
In [41]: # custm_df['SignupDate'] = pd.to_datetime(final_df['SignupDate'])
```

Merging all datasets into one based on their category

```
In [42]: combine_df = pd.merge(trans_df, product_df, on="ProductID", how="left")
```

```
In [43]: combine_df.head()
```

Out[43]:

	TransactionID	CustomerID	ProductID	TransactionDate	Quantity	TotalValue	Price_x	1
0	T00001	C0199	P067	2024-08-25 12:38:23	1	300.68	300.68	
1	T00112	C0146	P067	2024-05-27 22:23:54	1	300.68	300.68	
2	T00166	C0127	P067	2024-04-25 07:38:55	1	300.68	300.68	
3	T00272	C0087	P067	2024-03-26 22:55:37	2	601.36	300.68	
4	T00363	C0070	P067	2024-03-21 15:10:10	3	902.04	300.68	

In [44]: `final_df = pd.merge(combine_df, custm_df, on="CustomerID", how="left")`

In [45]: `final_df.head()`

Out[45]:

	TransactionID	CustomerID	ProductID	TransactionDate	Quantity	TotalValue	Price_x	1
0	T00001	C0199	P067	2024-08-25 12:38:23	1	300.68	300.68	
1	T00112	C0146	P067	2024-05-27 22:23:54	1	300.68	300.68	
2	T00166	C0127	P067	2024-04-25 07:38:55	1	300.68	300.68	
3	T00272	C0087	P067	2024-03-26 22:55:37	2	601.36	300.68	
4	T00363	C0070	P067	2024-03-21 15:10:10	3	902.04	300.68	

In [46]: `final_df.isnull().sum()`

```
Out[46]: TransactionID      0
         CustomerID      0
         ProductID       0
         TransactionDate  0
         Quantity        0
         TotalValue      0
         Price_x         0
         ProductName     0
         Category        0
         Price_y         0
         CustomerName    0
         Region          0
         SignupDate      0
         dtype: int64
```

```
In [47]: final_df.columns
```

```
Out[47]: Index(['TransactionID', 'CustomerID', 'ProductID', 'TransactionDate',
               'Quantity', 'TotalValue', 'Price_x', 'ProductName', 'Category',
               'Price_y', 'CustomerName', 'Region', 'SignupDate'],
              dtype='object')
```

```
In [48]: final_df.duplicated().sum()
```

```
Out[48]: 0
```

```
In [49]: final_df.dtypes
```

```
Out[49]: TransactionID      object
         CustomerID      object
         ProductID       object
         TransactionDate  object
         Quantity        int64
         TotalValue      float64
         Price_x         float64
         ProductName     object
         Category        object
         Price_y         float64
         CustomerName    object
         Region          object
         SignupDate      object
         dtype: object
```

Here signupDate and TransactionDate are in object so need to change it into date formate.

```
In [50]: final_df['TransactionDate'] = pd.to_datetime(final_df['TransactionDate'])
         final_df['SignupDate'] = pd.to_datetime(final_df['SignupDate'])
```

```
In [51]: final_df.dtypes
```

```
Out[51]: TransactionID      object
CustomerID      object
ProductID       object
TransactionDate  datetime64[ns]
Quantity        int64
TotalValue      float64
Price_x         float64
ProductName      object
Category        object
Price_y         float64
CustomerName     object
Region          object
SignupDate      datetime64[ns]
dtype: object
```

finding region wise customer distribution

```
In [52]: custm_dist = final_df['Region'].value_counts()
display(custm_dist)
```

```
Region
South America    304
North America    244
Europe           234
Asia             218
Name: count, dtype: int64
```

```
In [53]: # Here we can see high number of customer is distributed in South America followed .
# Marketing team has to give more efforts in rest of the region as the sellings are
```

```
In [54]: sns.countplot(data=final_df, y='Region', order=custm_dist.index, palette='Set2')

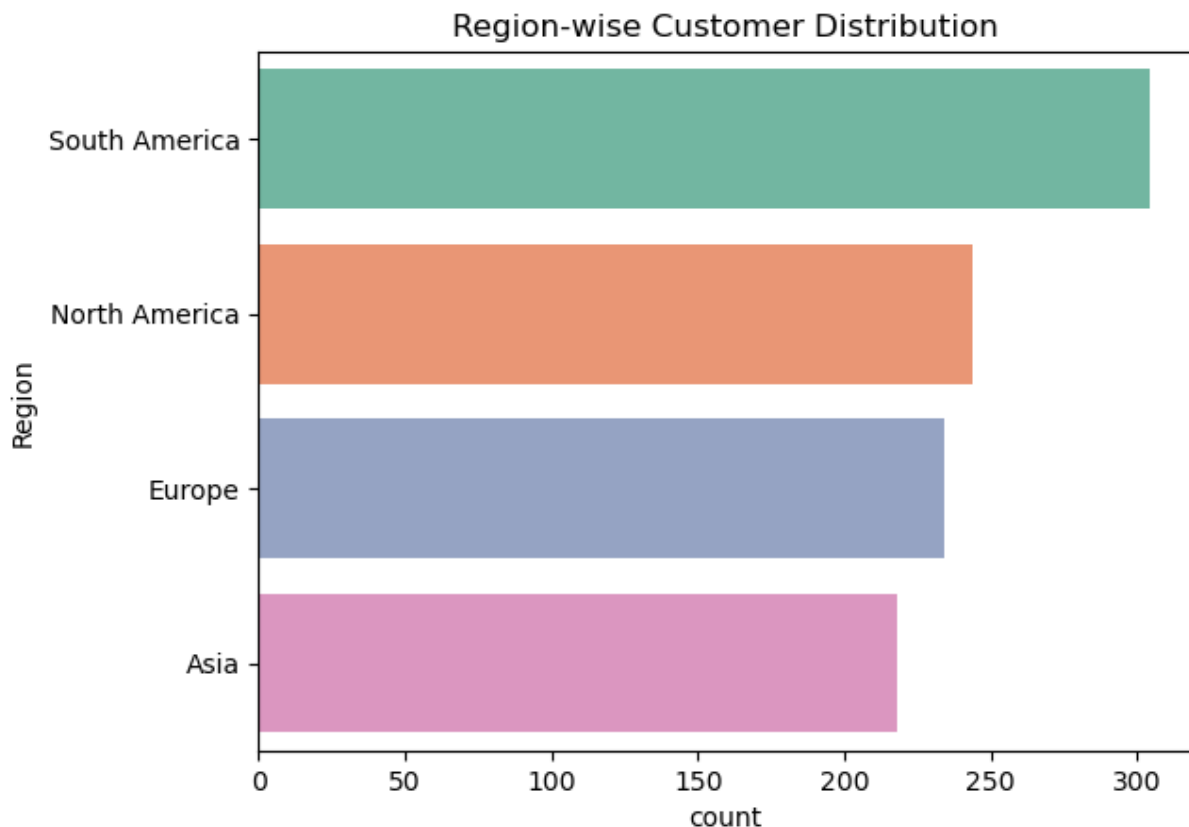
# Add a title
plt.title("Region-wise Customer Distribution")

# Show the plot
plt.show()
```

C:\Users\Lenovo\AppData\Local\Temp\ipykernel_2188\2608614170.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(data=final_df, y='Region', order=custm_dist.index, palette='Set2')
```



product analysis

```
In [55]: print("\nTop-selling Product Categories:")

top_product = final_df['Category'].value_counts()
display(top_product)
```

Top-selling Product Categories:

Category	
Books	270
Electronics	254
Home Decor	248
Clothing	228

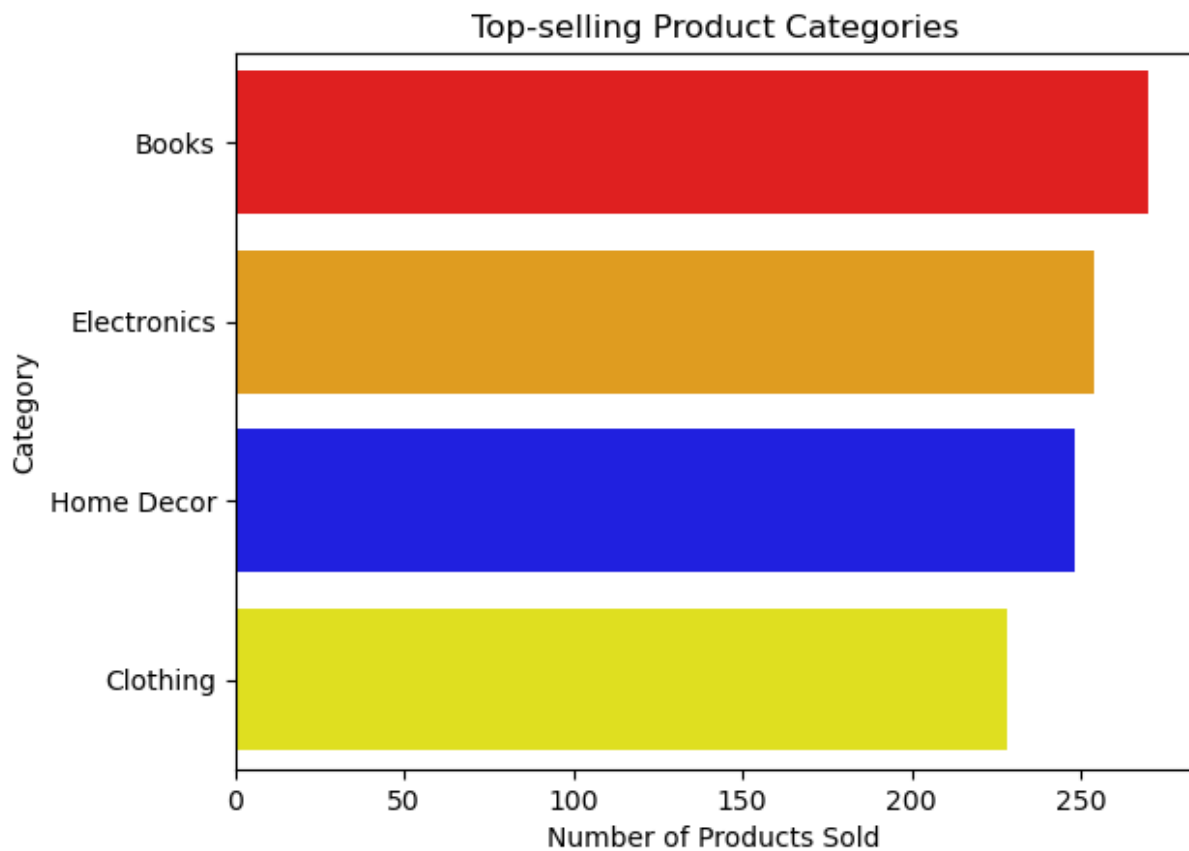
Name: count, dtype: int64

```
In [56]: # Visualize top categories
color = ('red', 'orange', 'blue', 'yellow')
sns.barplot(x=top_product.values, y=top_product.index, palette = color)
plt.title("Top-selling Product Categories")
plt.xlabel("Number of Products Sold")
plt.show()
```

C:\Users\Lenovo\AppData\Local\Temp\ipykernel_2188\2084253007.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=top_product.values, y=top_product.index, palette = color)
```



Transaction analysis

In [57]: *# Monthly revenue analysis*

```
final_df['Month'] = final_df['TransactionDate'].dt.to_period('M')
monthly_revenue = final_df.groupby('Month')['TotalValue'].sum()
monthly_revenue
```

Out[57]:

Month	
2023-12	3769.52
2024-01	66376.39
2024-02	51459.27
2024-03	47828.73
2024-04	57519.06
2024-05	64527.74
2024-06	48771.18
2024-07	71366.39
2024-08	63436.74
2024-09	70603.75
2024-10	47063.22
2024-11	38224.37
2024-12	59049.20


Freq: M, Name: TotalValue, dtype: float64

In [58]: *# Monthly revenue analysis via graph*

```
color = ('green')
monthly_revenue.plot(kind='line', figsize=(10, 5), marker='+', color = 'yellow', mar
plt.title("Monthly Revenue Trend")
```



```
plt.ylabel("Total Revenue")
plt.show()
```



finding region wise sales.

```
In [59]: # Total sales by region
total_sales = final_df.groupby('Region')['TotalValue'].sum().sort_values(ascending=
print("Total Sales by Region:\n", total_sales)
```

```
Total Sales by Region:
Region
South America    219352.56
Europe           166254.63
North America    152313.40
Asia             152074.97
Name: TotalValue, dtype: float64
```

```
In [60]: ## So as per the data we can clearly see South America generates the highest revenue
```

Best-Selling Products

```
In [61]: # Best-selling products by total sales
top_products = final_df.groupby('ProductName')['TotalValue'].sum().sort_values(ascending=
print("\nBest-Selling Products:\n", top_products)
```

```
Best-Selling Products:
ProductName
ActiveWear Smartwatch    39096.97
SoundWave Headphones     25211.64
SoundWave Novel          24507.90
ActiveWear Jacket        22712.56
ActiveWear Rug           22314.43
Name: TotalValue, dtype: float64
```

In [62]: *# The product 'ActiveWear Smartwatch' is the highest revenue generator, showing hig*

Active Regions by customers

In [63]: *# Count of customer as per activeness*

```
cust_active = final_df['Region'].value_counts()
print("\nMost Active Regions by customer:\n", cust_active)
```

Most Active Regions by customer:

Region	count
South America	304
North America	244
Europe	234
Asia	218

Name: count, dtype: int64

In [64]: *# South America have the highest active customer counts, suggesting more active cus*

Increasing Customer

In [66]: `custm_df['SignupDate'] = pd.to_datetime(final_df['SignupDate'])`

In [67]: *# Customer signups over the years*

```
signup_trends = custm_df['SignupDate'].dt.year.value_counts().sort_index()
print("\nCustomer Signup Trends:\n", signup_trends)
```

Customer Signup Trends:

SignupDate	count
2022	78
2023	56
2024	66


Name: count, dtype: int64

In [68]:

```
plt.figure(figsize=(8, 6))
sns.lineplot(x=signup_trends.index, y=signup_trends.values, palette="Blues_d")
plt.title("Customer Signup Trends Over the Years", fontsize=16)
plt.xlabel("Year", fontsize=14)
plt.ylabel("Number of Signups", fontsize=14)
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```

C:\Users\Lenovo\AppData\Local\Temp\ipykernel_2188\1394880443.py:2: UserWarning: Ignoring `palette` because no `hue` variable has been assigned.

```
sns.lineplot(x=signup_trends.index, y=signup_trends.values, palette="Blues_d")
```



In [69]: *# Here we can see the customer signup trend so need more focus to increase the tren*

Category wise selling trend

```
In [70]: # Total sales by product category
sales_by_category = final_df.groupby('Category')['TotalValue'].sum().sort_values(ascending=False)
print("\nSales by Product Category:\n", sales_by_category)
```

Sales by Product Category:

Category	TotalValue
Books	192147.47
Electronics	180783.50
Clothing	166170.66
Home Decor	150893.93

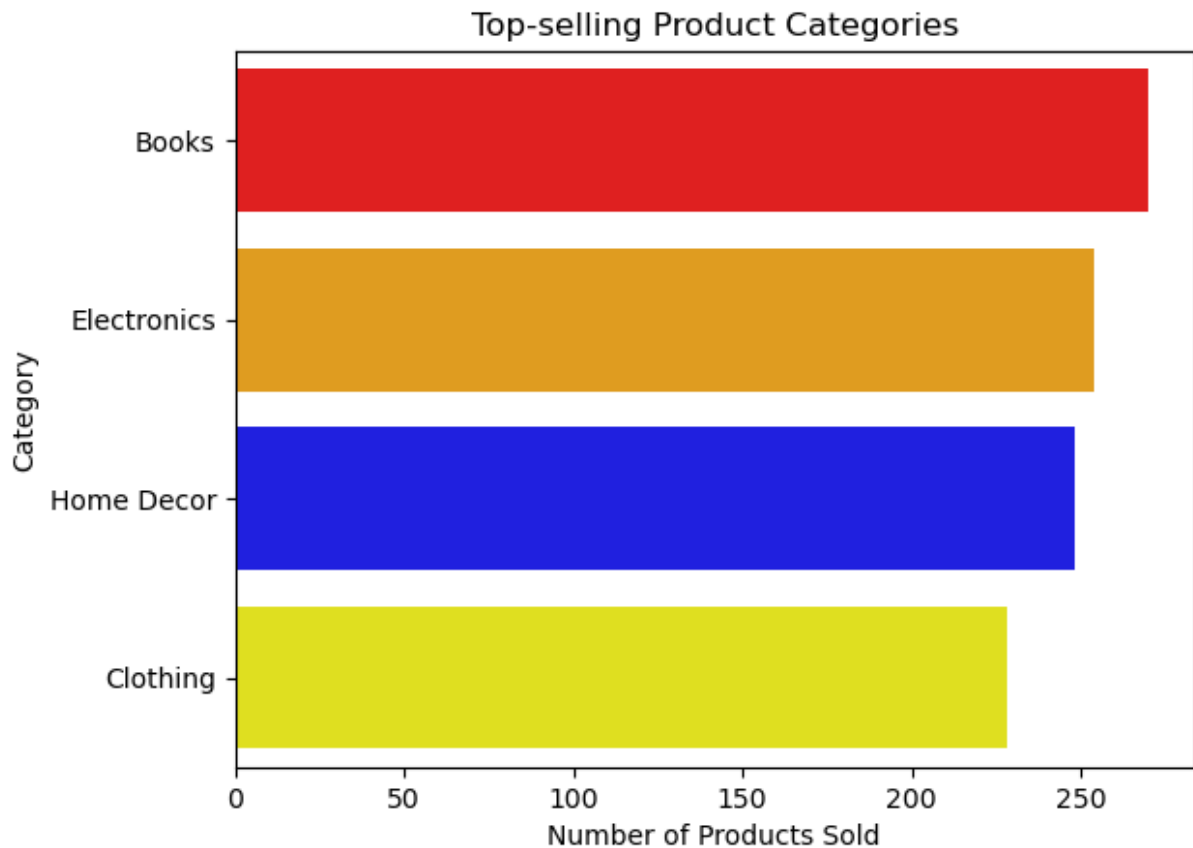
Name: TotalValue, dtype: float64

```
In [71]: color = ('red', 'orange', 'blue', 'yellow')
sns.barplot(x=top_product.values, y=top_product.index, palette = color)
plt.title("Top-selling Product Categories")
plt.xlabel("Number of Products Sold")
plt.show()
```

C:\Users\Lenovo\AppData\Local\Temp\ipykernel_2188\1540872222.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=top_product.values, y=top_product.index,palette = color)
```



The "Books" category drives the most revenue, highlighting that books remain a core product offering, followed by electronics.

Lookalike Model

```
In [87]: from sklearn.metrics.pairwise import cosine_similarity
from sklearn.preprocessing import StandardScaler
import pandas as pd

# Select relevant features for similarity
features = ['Region', 'TotalValue', 'Quantity']
numerical_data = final_df[features]

# Encode categorical data and scale numerical values
numerical_data = pd.get_dummies(numerical_data, columns=['Region'], drop_first=True)
```

```
scaler = StandardScaler()
scaled_data = scaler.fit_transform(numerical_data)

# similarity matrix
similarity_matrix = cosine_similarity(scaled_data)

# Create Lookalike recommendations for the first 20 customers
customer_ids = final_df['CustomerID'].unique()[:20]
lookalikes = {}

for customer in customer_ids:
    # Get the index of the current customer in the full dataset
    idx = final_df[final_df['CustomerID'] == customer].index[0]

    # Get similarity scores for the current customer
    scores = list(enumerate(similarity_matrix[idx]))

    # Sort scores in descending order of similarity
    scores = sorted(scores, key=lambda x: x[1], reverse=True)

    # Exclude the customer itself and get the top 3 similar customers
    recommendations = [
        (final_df.iloc[i]['CustomerID'], score)
        for i, score in scores
        if final_df.iloc[i]['CustomerID'] != customer
    ]
    lookalikes[customer] = recommendations[:3]

# Convert recommendations into a DataFrame
lookalike_df = pd.DataFrame([
    {
        "CustomerID": cust_id,
        "Lookalikes": [rec[0] for rec in recs],
        "Scores": [rec[1] for rec in recs]
    }
    for cust_id, recs in lookalikes.items()
])

# Print or save the results
print(lookalike_df)
```

	CustomerID	Lookalikes \
0	C0199	[C0127, C0051, C0009]
1	C0146	[C0145, C0028, C0161]
2	C0127	[C0199, C0051, C0009]
3	C0087	[C0130, C0155, C0126]
4	C0070	[C0051, C0127, C0041]
5	C0188	[C0158, C0187, C0163]
6	C0195	[C0087, C0003, C0004]
7	C0008	[C0029, C0156, C0047]
8	C0157	[C0024, C0156, C0016]
9	C0130	[C0087, C0155, C0126]
10	C0051	[C0070, C0127, C0041]
11	C0075	[C0062, C0066, C0017]
12	C0155	[C0087, C0130, C0126]
13	C0092	[C0088, C0002, C0173]
14	C0088	[C0092, C0002, C0173]
15	C0109	[C0059, C0049, C0018]
16	C0041	[C0170, C0073, C0086]
17	C0101	[C0054, C0028, C0092]
18	C0154	[C0049, C0008, C0072]
19	C0200	[C0084, C0145, C0110]

	Scores
0	[1.0000000000000002, 0.9999998401269015, 0.999...
1	[0.9999969792643577, 0.9999962162179905, 0.999...
2	[1.0000000000000002, 0.9999998401269015, 0.999...
3	[1.0000000000000002, 1.0000000000000002, 0.999...
4	[1.0, 0.9999977246008289, 0.9999977246008289]
5	[0.9999998114890177, 0.9999973292974776, 0.999...
6	[0.9999662462521426, 0.9985904179976799, 0.998...
7	[0.9999981433545493, 0.9999981433545493, 0.999...
8	[0.9999732972675586, 0.9999732972675586, 0.999...
9	[1.0000000000000002, 1.0000000000000002, 0.999...
10	[1.0, 0.9999977246008289, 0.9999977246008289]
11	[0.999987913428094, 0.9999848629040765, 0.9999...
12	[1.0000000000000002, 1.0000000000000002, 0.999...
13	[1.0, 0.9995338039385248, 0.9993519323400206]
14	[1.0, 0.9995338039385248, 0.9993519323400206]
15	[0.9998042773278407, 0.9991948075032455, 0.999...
16	[0.9999187104069156, 0.9998104989878519, 0.999...
17	[0.9994818808992558, 0.9948747032037053, 0.992...
18	[1.0, 0.9999720508484822, 0.9994082262803828]
19	[0.998481628265308, 0.9964429716059569, 0.9964...

```
In [93]: lookalike_df['Scores'] = [
          [rec[1] for rec in lookalikes[cust_id]] for cust_id in lookalike_df['CustomerID']
          ]
          print(lookalike_df.head())
```

	CustomerID	Lookalikes \
0	C0199	[C0127, C0051, C0009]
1	C0146	[C0145, C0028, C0161]
2	C0127	[C0199, C0051, C0009]
3	C0087	[C0130, C0155, C0126]
4	C0070	[C0051, C0127, C0041]

	Scores
0	[1.0000000000000002, 0.9999998401269015, 0.999...
1	[0.9999969792643577, 0.9999962162179905, 0.999...
2	[1.0000000000000002, 0.9999998401269015, 0.999...
3	[1.0000000000000002, 1.0000000000000002, 0.999...
4	[1.0, 0.9999977246008289, 0.9999977246008289]

```
In [97]: # Save to CSV
lookalike_df.to_csv('Lookalike.csv', index=False)
```

Customer segmentation , for clustering we will use k-means and evaluate the cluster using metrics

```
In [94]: from sklearn.cluster import KMeans
from sklearn.metrics import davies_bouldin_score

# Use numerical data for clustering
kmeans = KMeans(n_clusters=5, random_state=42)
clusters = kmeans.fit_predict(scaled_data)


# Add cluster Labels to the dataframe
final_df['Cluster'] = clusters

# Calculate Davies-Bouldin Index
db_index = davies_bouldin_score(scaled_data, clusters)
print(f"Davies-Bouldin Index: {db_index}")

# Visualize clusters
import seaborn as sns
import matplotlib.pyplot as plt

sns.scatterplot(x=scaled_data[:, 0], y=scaled_data[:, 1], hue=clusters, palette="vi")
plt.title("Customer Clusters")
plt.show()
```

Davies-Bouldin Index: 1.0015634918863825



In []:

In []:

In []:

In []: